

In [22]:

```
1 # Run this to ensure that you have the relevant packages:
2 import types
3 def imports():
4     for name, val in globals().items():
5         if isinstance(val, types.ModuleType):
6             yield val.__name__
7 list(imports())
```

Out[22]:

```
['builtins', 'builtins', 'numpy', 'pandas', 'matplotlib.pyplot', 'types']
```



In [1]:

```
1 import numpy as np
2 import pandas as pd
3 # I'll use pandas to manipulate data, but not to write the algo...
4
5 from IPython.core.interactiveshell import InteractiveShell
6 InteractiveShell.ast_node_interactivity = "all"
7
8
9 # loading the reacy, binarized data:
10 Xtraining = pd.read_csv('/Users/a8407352/Desktop/deepLearn/datasets/titanic/train.csv')
11 Xtraining = Xtraining.drop(Xtraining.columns[0],axis=1)
```

In [2]:

```

1 age_avg=Xtraining['0'].mean()
2 age_std=Xtraining['0'].std()
3 age_null_count =Xtraining['0'].isnull().sum()
4 age_null_random_list = np.random.randint(age_avg - age_std, age_avg + age_std, s
5 Xtraining['0'][np.isnan(Xtraining['0'])] = age_null_random_list
6 Xtraining['0'] = Xtraining['0'].astype(int)
7
8 Xtraining['CategoricalAge'] = pd.cut(Xtraining['0'], 5)
9
10 print (Xtraining[['CategoricalAge', '16']].groupby(['CategoricalAge'], as_index=
11
12 Xtraining.loc[ Xtraining['0'] <= 16, '0'] = 0
13 Xtraining.loc[(Xtraining['0'] > 16) & (Xtraining['0'] <= 32), '0'] = 1
14 Xtraining.loc[(Xtraining['0'] > 32) & (Xtraining['0'] <= 48), '0'] = 2
15 Xtraining.loc[(Xtraining['0'] > 48) & (Xtraining['0'] <= 64), '0'] = 3
16 Xtraining.loc[ Xtraining['0'] > 64, '0'] = 4
17

```

	CategoricalAge	16
0	(-0.08, 16.0]	0.201923
1	(16.0, 32.0]	0.229381
2	(32.0, 48.0]	0.172414
3	(48.0, 64.0]	0.246377
4	(64.0, 80.0]	0.181818

In [3]:

```

1 # Xtraining=Xtraining.drop(Xtraining['CategoricalAge'], inplace=True)
2 Xtraining = Xtraining.iloc[:,0:25]
3 Xtraining.head(10)

```

Out[3]:

	0	1	2	3	4	5	6	7	8	9	...	15	16	17	18	19	20	21	22	23	24
0	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0
1	2	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
2	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
3	2	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
4	2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0
5	2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0
6	3	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0
7	0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	1.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0
8	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
9	0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0

10 rows × 25 columns

In [4]:

```

1 yTrain = pd.read_csv('/Users/a8407352/Desktop/deepLearn/datasets/titanic/trainLa
2 yTrain = yTrain.drop(yTrain.columns[0],axis=1)
3 yTrain.head()

```

Out[4]:

	Survived
0	0
1	1
2	1
3	1
4	0

In [5]:

```

1 def sigmoid(x, deriv=False):
2     if deriv:
3         return x*(1-x)
4     return 1/(1+np.exp(-x))
5
6
7 weights1 = np.random.random((Xtraining.shape[1],1)) #bias can be added here, li
8 weights2 = np.random.random((1,1)) #L1.shape[0]

```

In [6]:

```

1 errors = []
2 index = []
3 for batch in range(10000):
4     L0 = Xtraining
5     L1 = sigmoid(Xtraining.dot(weights1)) #891X1
6     L2 = sigmoid(L1.dot(weights2))
7     # more layers here...possibly.
8     # error:
9     er = yTrain.values-L2 # 891X1
10    index.append(batch)
11    errors.append(np.mean(np.abs(er))) #since we take the mean of a batch of errors
12    #within the results, or smoothen results.
13    if batch%100==0:
14        print ('error:' +str(np.mean(np.abs(er))))
15
16
17
18 # learning / backprop:
19 # propogate:
20 # regular vector multiplication to find the delta:
21 delta1 = er*sigmoid(L2, deriv=True) # 2X1
22 er2 = delta1.dot(weights2.T) # 2X2
23 delta2 = er2 * sigmoid(L1, deriv=True) # 2X2
24 # update weights with gradient descent:
25 weights2 += L2.T.dot(delta1)
26 weights1 += L1.T.dot(delta2)
27

```

```

error:0    0.503945
dtype: float64
error:0    0.503955
dtype: float64

```

```

error:0    0.503965
dtype: float64
error:0    0.503974
dtype: float64
error:0    0.503983
dtype: float64
error:0    0.503992
dtype: float64
error:0    0.504001
dtype: float64
error:0    0.50401
dtype: float64
error:0    0.504018
dtype: float64
error:0    0.504026

```

In [7]:

```

1 # plotting:
2 import matplotlib.pyplot as plt
3 # print(errors)


```

In [8]:

```
1 plt.plot(index,errors)
```

Out[8]:

```
[<matplotlib.lines.Line2D at 0x116c0f2b0>]
```



In [9]:

```
1 weights1 = np.random.random((Xtraining.shape[1],1)) #bias can be added here, like weights1 = np.random.random((Xtraining.shape[1]+1,1))
2 weights2 = np.random.random((1,1)) #L1.shape[0]
```

In [10]:

```
1 def Relu(x, deriv=False):
2     if deriv:
3         x[x<=0] = 0
4         x[x>0] = 1
5         return x #notice that the derivative is not defined for x=0, but we'll take 0
6     else:
7         return np.maximum(0,x)
```

In [11]:

```
1 errors = []
2 index = []
3 for batchNum in range(100):
4     L0 = Xtraining
5     dotProduct = Xtraining.dot(weights1)
6     L1 = Relu(dotProduct) #891X1
7     L2 = Relu(L1.dot(weights2))
8
9     # obviously, we see by now that ReLu is not adequate for the task,
10    # i.e. for binary classification we need to squeeze the output into a 0 to 1 range
11    # or in other words, probability.
```

In [12]:

```
1 # Let's change the loss function - perhaps a log loss would be better.
2 def cross_entropy(predictions, targets):
3     N = predictions.shape[0]
4     ce = -np.sum(targets*np.log(predictions))/N
5     return ce
```

```

1 errors = []
2 index = []
3 for batch in range(10000):
4     L0 = Xtraining
5     L1 = sigmoid(Xtraining.dot(weights1)) #891X1
6     L2 = sigmoid(L1.dot(weights2))
7     # more layers here...possibly.
8     # error:
9     er = cross_entropy(L2,yTrain.values) # 891X1
10    index.append(batch)
11    errors.append(np.mean(np.abs(er)))
12    if batch%100==0:
13        print ( 'error:' +str(np.mean(np.abs(er))))
14
15
16
17 # learning / backprop:
18 # propagate:
19 # regular vector multiplication to find the delta:
20    delta1 = er*sigmoid(L2, deriv=True) # 2X1
21    er2 = delta1.dot(weights2.T) # 2X2
22    delta2 = er2 * sigmoid(L1, deriv=True) # 2X2
23 # update weights with gradient descent:
24    weights2 += L2.T.dot(delta1)
25    weights1 += L1.T.dot(delta2)
26

```

[illegible]

[illegible]



```
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
error:7.889626101711672e-13
```

In [14]:

```
1 plt.plot(index,errors)
```

Out[14]:

[<matplotlib.lines.Line2D at 0x114ac99b0>]



In [15]:

```
1 # Not much different (the error is calculated differently, but the learning rate
2 # not getting much better). I tend to think that this is more because of the arch
3 # which is quite arbitrary. The network takes the form of 891X891 two layers.
4 # I will plot a similar one with 5X5 two layers:
```

In [16]:

```

1  import matplotlib.pyplot as plt
2
3  def draw_neural_net(ax, left, right, bottom, top, layer_sizes):
4      '''
5          Draw a neural network cartoon using matplotlib.
6
7          :usage:
8              >>> fig = plt.figure(figsize=(12, 12))
9              >>> draw_neural_net(fig.gca(), .1, .9, .1, .9, [4, 7, 2])
10
11          :parameters:
12              - ax : matplotlib.axes.AxesSubplot
13                  The axes on which to plot the cartoon (get e.g. by plt.gca())
14              - left : float
15                  The center of the leftmost node(s) will be placed here
16              - right : float
17                  The center of the rightmost node(s) will be placed here
18              - bottom : float
19                  The center of the bottommost node(s) will be placed here
20              - top : float
21                  The center of the topmost node(s) will be placed here
22              - layer_sizes : list of int
23                  List of layer sizes, including input and output dimensionality
24      '''
25      n_layers = len(layer_sizes)
26      v_spacing = (top - bottom)/float(max(layer_sizes))
27      h_spacing = (right - left)/float(len(layer_sizes) - 1)
28      # Nodes
29      for n, layer_size in enumerate(layer_sizes):
30          layer_top = v_spacing*(layer_size - 1)/2. + (top + bottom)/2.
31          for m in range(layer_size):
32              circle = plt.Circle((n*h_spacing + left, layer_top - m*v_spacing), v
33                                  color='w', ec='k', zorder=4)
34              ax.add_artist(circle)
35      # Edges
36      for n, (layer_size_a, layer_size_b) in enumerate(zip(layer_sizes[:-1], layer
37          layer_top_a = v_spacing*(layer_size_a - 1)/2. + (top + bottom)/2.
38          layer_top_b = v_spacing*(layer_size_b - 1)/2. + (top + bottom)/2.
39          for m in range(layer_size_a):
40              for o in range(layer_size_b):
41                  line = plt.Line2D([n*h_spacing + left, (n + 1)*h_spacing + left
42                                     [layer_top_a - m*v_spacing, layer_top_b - o*v
43                  ax.add_artist(line)
44


```

In [17]:

```
1 fig = plt.figure(figsize=(12, 12))
2 ax = fig.gca()
3 ax.axis('off')
4 draw_neural_net(ax, .1, .9, .1, .9, [6, 6])
5 fig.savefig('nn.png')
6
```

Out[17]:

(0.0, 1.0, 0.0, 1.0)



In [18]:

```
1 # As can be seen, this doesn't really make a whole lot sense and further experim
```

In [20]:

```
1 # References:  
2 # I used numerous tutorials and stackoverflow samples whose code I adopted, put  
3 # to have what I tried to achieve as a playground for the most simple ANNs.
```